


NN: Model interpretability: LIME

LIME


How to make interpretability for a Complex model like NN?

What approach can be used for model interpretability?




- Lime (Local Interpretable Model-agnostic Explanations) is used to explain predictions made by machine learning models.
- Lime is a popular technique for interpreting individual predictions of black-box models.
- It approximates the local behavior of a model by creating a local interpretable model, which helps us understand why a model makes a specific prediction for a given instance.

Understanding Global Analysis



Understanding Local Analysis

Now, instead of entire dataset
 we look at a particular datapoint



In local Analysis:

↳ finding →

what is special about \bar{x}_q point & how varying it causes model predictions to change

↳ called as Local Interpretability

Understanding locality of interest

How to vary \bar{x}_q datapoint?

↳ for this we will create N datapoints which will be neighbors to \bar{x}_q .

$$\therefore \bar{X} = \left\{ \bar{x}_q + \left(\epsilon_i \right) \right\}_{i=1}^{N=10,000}$$

→ Adds noise to our \bar{x}_q datapoint

& with $\|z_i\|^2 \leq \alpha$



this means that all the datapoints will lie inside a circular region with radius = α

creates locality of interest around x_q

Note: The process is analogous to Nearest Neighbor

These Locality of interest points shows how fluctuating features

- Affects the model's performance

Understanding model Agnostics

③ Understanding Model Agnostic

Remember for

* linear model \Rightarrow wts are checked for interpretability

* Decision model \Rightarrow Depth of the tree is used for interpretability

Observe:

* for different models, different interpretability techniques used

Qs LIME used only for NN?

\hookrightarrow No, LIME is a general interpretability technique

* Hence can be used for any model

\Downarrow
therefore LIME is model Agnostic

What has LIME had to offer on model interpretability?

1. A consistent model agnostic explainer [LIME].
2. A method to select a representative set with explanations [SP-LIME] to make sure the model behaves consistently while replicating human logic. This representative set would provide an intuitive global understanding of the model.


LIME explains a prediction so that even the non-experts could compare and improve on an untrustworthy model through feature engineering. An ideal model explainer should contain the following desirable properties:

Summarizing LIME process

Summarizing steps of LIME:

Step 1: select a local point \bar{x}_q

Step 2: create N datapoints within α radius of \bar{x}_q by introducing noise to the local point



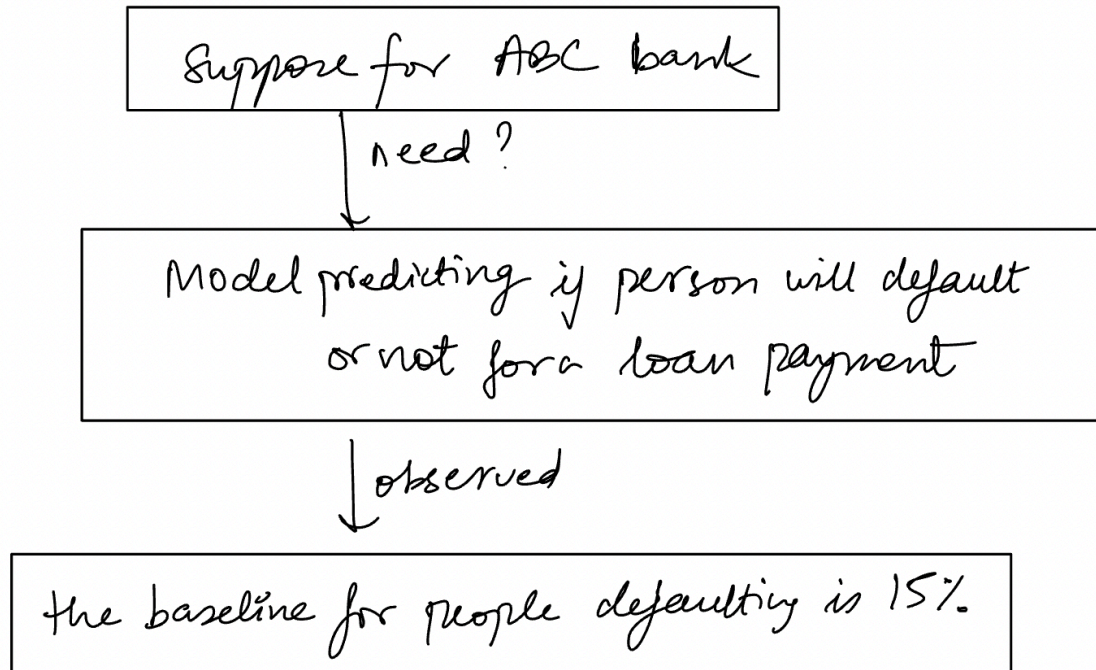
Step 3: Pass all the points in the new dataset to the model

Step 4: Make a new interpretable (simpler) model to learn on the newly created dataset

Note: This newly created interpretable model uses k -Lasso as loss function.
↳ finds top k features using $L1$ regularization


SHAP: Alternate way to find interpretability

SHAP: Another way to interpret model



Suppose for a person B, the model predicts 25%,
then why model prediction for person B differs
from our baseline?

↳ this is what SHAP measures



- SHAP values (SHapley Additive exPlanations) is a method rooted in cooperative game theory that enhances the transparency and interpretability of machine learning models.
- It provides a unified framework for attributing the contributions of each feature towards a prediction, addressing the limitations of traditional linear models and feature importance in tree-based models.
- Unlike linear models that rely on feature coefficients, which may be influenced by variable scales and fail to capture local importance, SHAP values offer a more comprehensive understanding of feature importance.
- They consider the impact of each feature when combined with different subsets of features, accounting for interactions and dependencies among them.
- SHAP values are based on the concept of Shapley values from cooperative game theory. By calculating the marginal contribution of a feature when added to or removed from coalitions of features, SHAP values ensure fairness and consistency in attributing importance.
- The insights provided by SHAP values are interpretable and nuanced, revealing the relative importance of each feature in a prediction. Visualizations, such as summary plots, individual feature importance plots, dependence plots, or force plots, allow for a detailed exploration of feature impact on predictions.