

Tensorflow and Keras 1

Tensorflow

TensorFlow is the premier open-source deep learning framework developed and maintained by Google

Keras

- Using TensorFlow directly can be challenging,
- In TensorFlow 2, Keras has become the default high-level API
- No need to separately install keras
- The modern tf.keras API brings Keras's simplicity and ease of use to the TensorFlow project.

Why keras?

- easy to build and use giant deep learning models
- **Light-weight and quick**
- can support other backends as well besides tensorflow, eg: Theano
- Open source

Ways of writing code in Keras

- Sequential API
- Functional API

Keras Sequential API

- The simplest and recommended API to start with
- Called as "Sequential" because we add layers to the model one by one in a
- linear manner, from input to output.
- You can select optimizers, loss functions, and metrics while writing code

The dense layer helps us define one layer of a Feedforward NN.

Example model using Sequential API

```
[ ] model = Sequential([
    Dense(64, activation="relu"), #hidden dense layer with 64 neuron units
    Dense(4, activation="softmax") #output layer with 4 units and softmax activation
])
```

Note: The layers in the sequential model interact with each other therefore we don't need to define the input shape for all the layers.

We can check model weights using `model.weights`

model.add()

- Instead of passing the list of layers as an argument while creating a model instance, we can use the add method.

```
▶ model = Sequential()
model.add(Dense(64, activation="relu", input_shape=(11,)))
model.add(Dense(4, activation="softmax"))
```

model.summary()

To print the summary of the model we have created

```
[ ] model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 64)	768
dense_5 (Dense)	(None, 4)	260

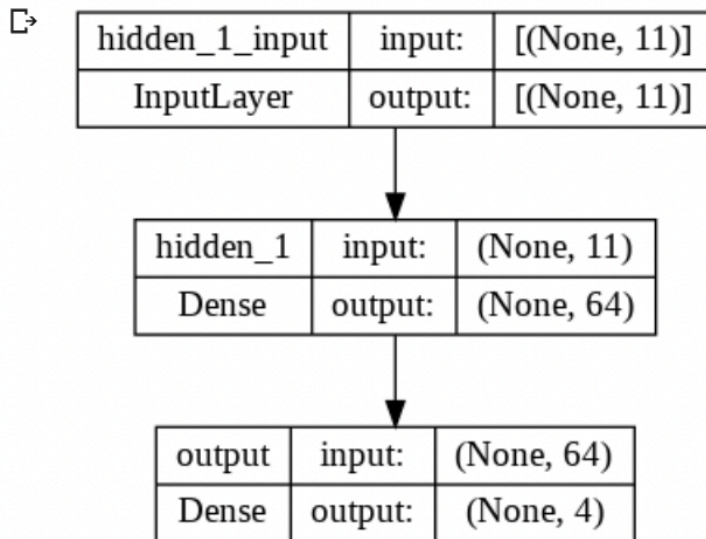
=====

Total params: 1,028
Trainable params: 1,028
Non-trainable params: 0

Plotting model

```
from tensorflow.keras.utils import plot_model

plot_model(model,
            to_file='model.png',
            show_shapes=True, show_layer_names=True)
```



Weights and Bias Initializer

in the Dense layer,

- 1) the biases are set to zero (zeros) by default
- 2) the weights are set according to glorot_uniform

For our own custom initializer, we can use **bias_initializer** and **kernel_initializer**

Compiler - loss and optimizer

What things to decide while compiling the model?

- Loss
- Optimizer

we can define a list of metrics that we might want to track during the training, like accuracy

```
[ ] model_2C = Sequential([
    Dense(64, activation="relu", input_shape=(11,)),
    Dense(1, activation="sigmoid")])

# new piece of code
model_2C.compile(
    optimizer = "adam", # stochastic gradient descent, adam, rmsprop, adadelts
    loss = "binary_crossentropy", # sigmoid loss, # mean_squared_error, categorical_crossentropy,
                                     #sparse_categorical_crossentropy, binary_crossentropy
    metrics = ["accuracy"]
)
```

- Another way to define **optimizer = keras.optimizers.Adam(learning_rate=0.01)**
- We can also pass customized loss and optimizer functions in keras models.
- These metrics will be calculated and saved after each epoch (one pass of whole data to update the model).

Epoch

- To avoid memory issues, data is passed in small batches instead of whole
- Each pass of the mini-batch is called an iteration.
- Each pass of whole datasets is called an Epoch.
- One epoch means that each sample in the training dataset has had an opportunity to update the internal model parameters.

Training: model.fit ()

It means updating the weights using the optimizer and loss functions on the dataset.

```
model.fit(X_train, y_train)
```

```
X_train = (num_samples, num_features)
```

```
y_train = (num_samples, num_classes) or y_train = (num_samples, )
```

Arguments we can pass to this method

- epochs (number of epochs you want to train for)
- batch_size (Batch size usually in form of 2^x like 4,8,16,32)
- Validation_split (size of validation data)
- verbose (0 for silent training, 1 to print each iteration)

```

%%time
model.fit(X_train, y_train, epochs=10, batch_size=256, validation_split=0.1, verbose=1)

## no of iterations: ( 10847 (training size) - 1084.7 (validation split) )/(256) == 39

Epoch 1/10
39/39 [=====] - 1s 8ms/step - loss: 1.3526 - accuracy: 0.3514 - val_loss: 1.2917 - val_accuracy: 0.4184
Epoch 2/10
39/39 [=====] - 0s 3ms/step - loss: 1.2592 - accuracy: 0.4327 - val_loss: 1.2054 - val_accuracy: 0.4525
Epoch 3/10
39/39 [=====] - 0s 3ms/step - loss: 1.1722 - accuracy: 0.4666 - val_loss: 1.1224 - val_accuracy: 0.4876
Epoch 4/10
39/39 [=====] - 0s 3ms/step - loss: 1.0951 - accuracy: 0.4942 - val_loss: 1.0554 - val_accuracy: 0.5051
Epoch 5/10
39/39 [=====] - 0s 3ms/step - loss: 1.0361 - accuracy: 0.5233 - val_loss: 1.0032 - val_accuracy: 0.5318

```

Note: After training, weights now will follow normal distribution and biases will not be Zero now

History object

- `model.fit` returns a history object which contains the record of progress NN training.
- History object contains records of loss and metrics values for each epoch.
- It's an alternative to `dir()`. `__dict__` attribute can be used to retrieve all the keys associated with the object on which it is called.

```

history = model.fit(X_train, y_train, validation_data = (X_val, y_val), epochs=500, batch_size=512, verbose=0)
history.__dict__.keys()

dict_keys(['validation_data', 'model', '_chief_worker_only', '_supports_tf_logs', 'history', 'params', 'epoch'])

```

The model has saved all the loss and metrics values for each epoch inside the history dictionary where all the values are stored in different lists.

```

epochs = history.epoch
loss = history.history["loss"]
accuracy = history.history["accuracy"]
val_loss = history.history["val_loss"]
val_accuracy = history.history["val_accuracy"]

```

Prediction and Evaluation

Evaluate the model

```
loss, accuracy = model.evaluate(X_test, y_test)
```

- `model.evaluate` returns the loss value & metrics value for the model.

- **weights/parameters are not updated during evaluation** (and prediction) which means **only forward pass, no backward pass**

```
[ ] loss, accuracy = model.evaluate(X_test, y_test)
    print('Test Set')
    print("Loss value : ", loss)
    print("Accuracy   : ", accuracy)
```

```
42/42 [=====] - 0s 2ms/step - loss: 0.5842 - accuracy: 0.7634
Test Set
Loss value : 0.5841851830482483
Accuracy   : 0.7634328603744507
```

Predictions

pred = model.predict(X_test)

```
[ ] pred = model.predict(X_test)
    pred
```

```
42/42 [=====] - 0s 1ms/step
array([[9.99999940e-01, 2.10884883e-17, 1.79274864e-33, 0.00000000e+00],
       [1.13163900e-03, 1.37660122e-02, 1.01408757e-01, 8.83693516e-01],
       [1.93726644e-02, 2.39155009e-01, 2.92279452e-01, 4.49192822e-01],
       ...,
       [1.64477840e-01, 8.35509479e-01, 1.26833165e-05, 9.64229950e-15],
       [6.74094200e-01, 3.25786144e-01, 1.19630786e-04, 1.66241088e-12],
       [1.01807564e-02, 2.50488132e-01, 7.39329159e-01, 2.02724095e-06]],
      dtype=float32)
```

- To get predictions on unseen data, **model.predict** method is used
- It returns the raw output from the model (i.e. probabilities of an observation belonging to each one of the 4 classes)
- the sum of probabilities of an observation belonging to each of the 4 classes will be 1 i.e., **np.sum(pred, axis=1)**

To know the class an observation belongs to, using these 4 probability values

- Find the index having the largest probability and that will be the predicted class.
- **pred_class = np.argmax(pred, axis = 1)**

To check the accuracy of the model using sklearn's **accuracy_score**

```
from sklearn.metrics import accuracy_score
acc_score = accuracy_score(y_test, pred_class)
```